

Mining Massive Datasets

Robert T. Bauer

Why Massive Datasets?

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer

- Having the data creates opportunities to improve the world
- Having the data creates entry barriers.
- Having the data allows you to make an inference others can't – this makes entry nearly impossible.

At a very fundamental level, we advance knowledge in two ways:

- ① Systems of axioms as sources of new theories (called autonomous systems of axioms)
- ② Given a set B , usually large and possibly infinite, of preassigned propositions, find a system of axioms from which the propositions follow. This gives the “theory” a physical (more generally, an empirical) basis.

Massive datasets provide an empirical foundation for a heterogeneous system of axioms.

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer



Map-Reduce Computation Schematic I

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication


Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer



Map-Reduce Computation Schematic II

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

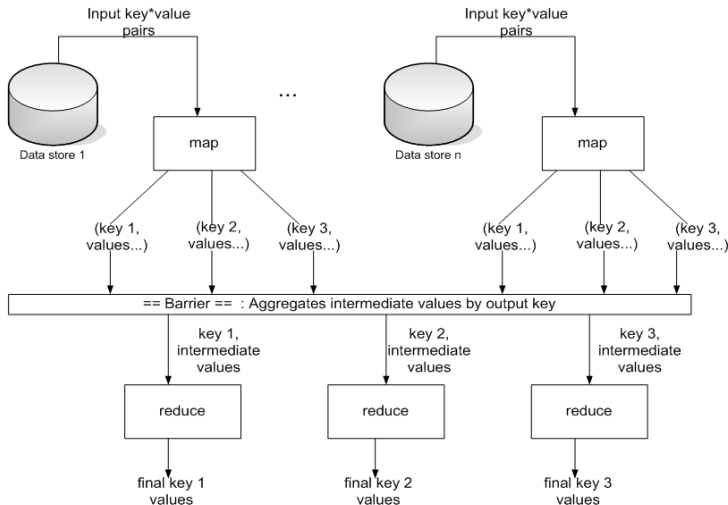
Relational Algebra

Conclusion


Closing Thoughts

Discussion

©2013 R.T.Bauer



Map-Reduce Key-Value Pair



Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer

Map-Reduce Execution

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

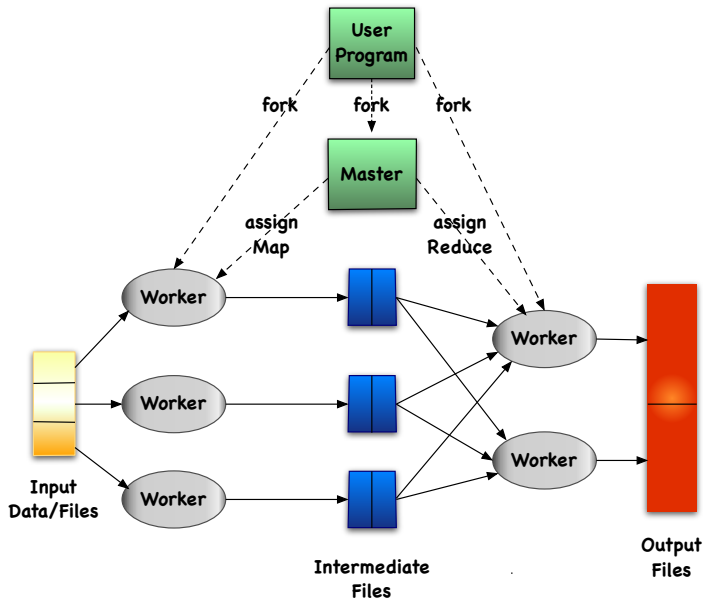
Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer



MR isn't "the" solution - nor is it the solution for every problem requiring parallel compute nodes. MR works well when we have

very large files that are rarely updated in place.

For example,

- No: Responding to product requests
- Yes: Large matrix-vector multiplication
- No: Recording sales
- Yes: Finding users whose buying patterns are similar
- Yes: Counting

MR Example: Counting Words I

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer

```
text_files=glob.glob('MyFiles/*.txt')

def getFileContents(fname):
    return open(fname).read()

source = dict((fn, getFileContents(fn))
              for fn in text_files)

# tell what do with the MR result
fout = open('MyFiles/outCount','w')
def final(k,v):
    print k,v
    f.write(str((k,v)))
```

MR Example: Counting Words II

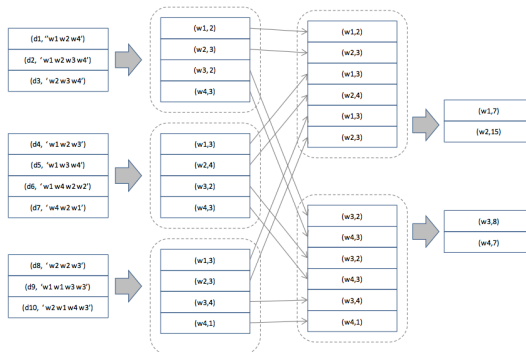
```
# the map function
#  ‘‘return’’ pair (word,1), for each word
def mapfn(k,v):
    for l in v.splitlines():
        for w in l.split():
            yield w.lower(), 1

# the reduce function
#  compute the ‘‘count’’
def reducefn(k,v)
    return k, len(v)
```

MR Example: Counting Words III

Map: $(d_k, 'w_1 \dots w'_n) \rightarrow [(w_i, c_i)]$ document -> word-count pairs

Reduce: $(w_i, [c_i]) \rightarrow (w_i, \sum_i c_i)$ word, count-list -> word-count-total



It's a math problem to determine the optimal number of mappers and reducers; not a tuning problem. Why?

Say we have an “internet” with 3 websites (s_0 , s_1 & s_2) with the following probabilities of leaving one site for another:

$$P = \begin{bmatrix} 0.9 & 0.075 & 0.025 \\ 0.15 & 0.8 & 0.05 \\ 0.25 & 0.25 & 0.5 \end{bmatrix}$$


The ergodic behaviour, usually called the *stationary* distribution (“visitation” ratio in performance modelling):

$$\lim_{n \rightarrow \infty} P^n = \begin{bmatrix} 0.625 & 0.3125 & 0.0625 \\ 0.625 & 0.3125 & 0.0625 \\ 0.625 & 0.3125 & 0.0625 \end{bmatrix}$$

tells us that 62.5% of the time, site s_0 is “visited”, 31.25% it is site s_1 and 6.25% of the time it is site s_2 — site s_0 is the most *popular* and site s_2 , the least.

[Introduction](#)[Map-Reduce](#)[Schema I](#)[Schema II](#)[Key-Value Pairs](#)[Execution](#)[Algorithms Using](#)[Counting](#)[Multiplication](#)[Relational Algebra](#)[Conclusion](#)[Closing Thoughts](#)[Discussion](#)[©2013 R.T.Bauer](#)

Matrix Multiplication



Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer

Multiplying (large) matrices is fundamental to big data.

$$x_i = \sum_j^n m_{ij} v_j$$

When n is on the order of 10^2 map-reduce or even a distributed file system (DFS) isn't the right choice. But what to do when n is on the order of 10^{10} ?

Let's tackle the case where n is large, but everything fits into memory.

Map-Reduce (fits in memory): Matrix \times Vector

We store M and v in memory.

We want the reduce function to get key-value pairs:

$$(i, m_{ij}v_j)$$

The reduce function sums these pairs, giving (i, x_i) .

Each map task is designed to handle one “chunk” of M . For example, given k map tasks, let


$$\text{mti} = ((i - 1)n + j) \bmod k$$

be the map task id for processing matrix element m_{ij} . The “source” is then the key-value pair $(\text{mti}, (i, j))$.

M and v are (in memory and) available to each map task; the map function, as mentioned above, emits $(i, m_{ij}v_j)$.

Map-Reduce (mem too small): Matrix \times Vector

The technique is to break M into square chunks and to partition v into compatible regions.



A (relational) *links* table:

From	To
ulr1	url2
url1	url3
url2	url3
url2	url4
...	...

(Some) Relational Operations:

Selection: “where” clause; $\sigma_C(R)$

Projection: “select” attributes; $\pi_S(R)$

Join: “natural” - match common attributes; $T \bowtie U$

[Introduction](#)[Map-Reduce](#)[Schema I](#)[Schema II](#)[Key-Value Pairs](#)[Execution](#)[Algorithms Using](#)[Counting](#)[Multiplication](#)[Relational Algebra](#)[Conclusion](#)[Closing Thoughts](#)[Discussion](#)[©2013 R.T.Bauer](#)

Selection, $\sigma_C(R)$, is very simple and can be done either in the map or reduce function.

Map: For each tuple $t \in R$, if t satisfies C emit key-value pair (t, t) .

Reduce: Emit each key-value pair; reduce is the “identity” function.

Projection, $\pi_S(R)$, filters out attributes not in S - this can cause duplicate “rows” in the output.

We use the reduce function to eliminate duplicates.

Map: For each tuple $t \in R$, construct tuple t' by removing components not in the projection.

Emit the key-value pair (t', t') .

Reduce: The map-reduce aggregation

$$[(t', t')] \rightarrow (t', [t', t', \dots, t'])$$

passes $(t', [t', t', \dots, t'])$ to the reduce function which emits (t', t') .

Map-Reduce - Natural Join (Map)

We illustrate a natural join by finding all 2 hop links in the “from-to” table. For example, if we have $\delta(u, t)$ and $\delta(t, v)$, then we can determine that v is “reachable” from u in two hops (i.e. $\delta(u, t) \wedge \delta(t, v) \Rightarrow \delta(u, v)$).

Map: For each tuple $(u, t) \in \text{from-to}$ emit two key-value pairs:

$$(t, (\mathbf{ONE}, u)) \quad (u, (\mathbf{TWO}, t))$$

Consider the *links*:

$$\delta(u, t) \quad \delta(w, t) \quad \delta(t, x) \quad \delta(t, v)$$

Map emits:

$$\begin{array}{ll} (t, (\mathbf{ONE}, u)) & (u, (\mathbf{TWO}, t)) \\ (t, (\mathbf{ONE}, w)) & (w, (\mathbf{TWO}, t)) \\ (x, (\mathbf{ONE}, t)) & (t, (\mathbf{TWO}, x)) \\ (v, (\mathbf{ONE}, t)) & (t, (\mathbf{TWO}, v)) \end{array}$$

Map-Reduce - Natural Join (Reduce)

The map-reduce aggregation combines:

$$\begin{array}{ll} (t, (\mathbf{ONE}, u)) & (u, (\mathbf{TWO}, t)) \\ (t, (\mathbf{ONE}, w)) & (w, (\mathbf{TWO}, t)) \\ (x, (\mathbf{ONE}, t)) & (t, (\mathbf{TWO}, x)) \\ (v, (\mathbf{ONE}, t)) & (t, (\mathbf{TWO}, v)) \end{array}$$

To get:

$$\begin{array}{ll} (x, [(\mathbf{ONE}, t)]) & (v, [(\mathbf{ONE}, t)]) \\ (u, [(\mathbf{TWO}, t)]) & (w, [(\mathbf{TWO}, t)]) \end{array}$$

$$(t, [(\mathbf{ONE}, u), (\mathbf{ONE}, w), (\mathbf{TWO}, x), (\mathbf{TWO}, v)])$$

Reduce: Cartesian product ($\mathbf{ONE} \times \mathbf{TWO}$) and filter out the intermediate (t) hop, emitting:

$$(u, x), (u, v), (w, x), (w, v)$$

Map-Reduce SQL Query I

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

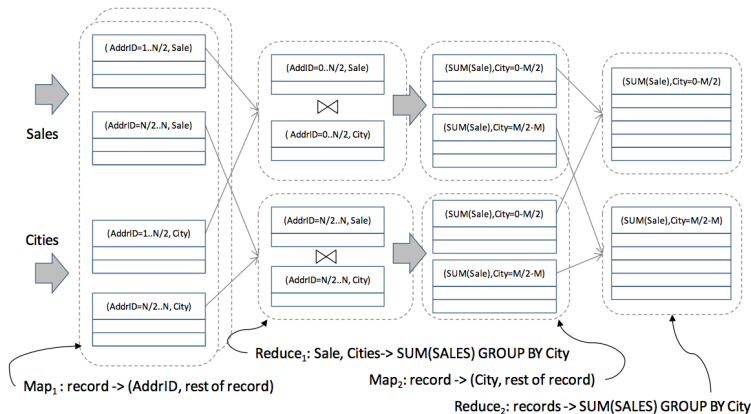
Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer



SQL: SELECT SUM(Sale), City FROM Sales, Cities WHERE Sales.AddrID=Cities.AddrID GROUP BY City

- ① **MAP**₁ emits (addrId, sale) or (addrId, city)
depending on “type”
- ② **AGGREGATION**₁ gives (addrId, [city, sale₁, ...])
- ③ **REDUCE**₁ emits (city, $\sum \text{sale}_i$)
- ④ **AGGREGATION**₂ gives (city, [$\sum \text{sale}$, $\sum \text{sale}$, ...])
- ⑤ **REDUCE**₂ emits (city, $\sum \sum \text{sale}$)

Map-Reduce SQL Query III

Introduction

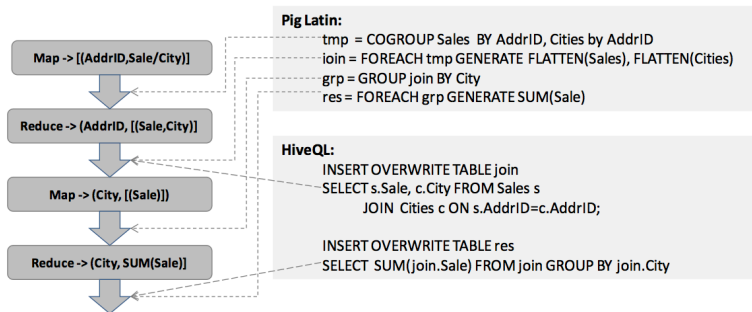
Map-Reduce

Schema I
Schema II
Key-Value Pairs
Execution
Algorithms Using
Counting
Multiplication
Relational Algebra

Conclusion

Closing Thoughts
Discussion

©2013 R.T.Bauer



SQL: SELECT SUM(Sale), City from Sales, Cities WHERE Sales.AddrID=Cities.AddrID GROUP BY City

Map-Reduce is fairly straight forward. Understanding the aggregation step is key; there are lots of patterns.

In January, I will cover computational complexity of map-reduce. This will allow us to determine optimal layout of data, how many mappers, reducers, and best way to aggregate data. Working with the computational models also allows us to determine whether map-reduce even makes sense for the problem (there are other parallel approaches).

I hope you have a better understanding of map-reduce and the basic patterns. You can download octo or mincemeat and play with map-reduce on your desktop.

Next month, I will go over recommender systems.

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer

Q & A

Introduction

Map-Reduce

Schema I

Schema II

Key-Value Pairs

Execution

Algorithms Using

Counting

Multiplication

Relational Algebra

Conclusion

Closing Thoughts

Discussion

©2013 R.T.Bauer